

Advanced Face Recognition System for Enhanced Image Analysis

Miraj Kashyap Vyas
(1260233)
School of Engineering
University of Guelph
mirajkas@uoguelph.ca

Harshitsinh Chauhan
(1288639)
School of Engineering
University of Guelph
harshits@uoguelph.ca

Abstract— Face recognition technology has witnessed remarkable advancements in recent years, revolutionizing various sectors including security, surveillance, and biometrics. This report provides a comprehensive overview of the latest developments in image analysis techniques for face recognition. The goal of this project is to create a sophisticated Face Recognition System (FRS) that is suited for demanding image analysis applications. By utilizing deep learning methods, specifically Convolutional Neural Networks (CNNs), the system seeks to improve face recognition accuracy and dependability in a variety of contexts. This report explores the integration of Convolutional Neural Networks (CNN) and Histogram of Oriented Gradients (HOG) for face detection. Leveraging CNN's ability to automatically learn features and HOG's strength in capturing local texture information, the study investigates their combined efficacy in accurately detecting faces within images. Through a comprehensive analysis, the report evaluates the performance of the integrated approach against traditional methods. Experimental results showcase the effectiveness of CNN-HOG fusion in detecting faces across diverse datasets, highlighting its robustness and scalability. Additionally, the report delves into the computational efficiency and practical implications of deploying this hybrid model in real-world scenarios. Insights from the study contribute to advancing the field of computer vision, offering valuable implications for applications such as surveillance, security, and human-computer interaction

Keywords—Face Recognition, CNN, HOG, Detection

I. INTRODUCTION

Face identification is an essential part of many computer vision applications, such as automatic social networking platform tagging, security systems, human-computer interface, and surveillance. Histogram of Oriented Gradients (HOG) algorithms and Convolutional Neural Networks (CNN) have gained recognition recently as powerful methods for handling face identification challenges. [1]

In this study, the face detection and face recognition modules are used to explore the combination of CNN and HOG algorithms. The main goal is to use the advantages of both approaches to improve face detection performance in image analysis. Through the integration of CNNs' strong feature learning capabilities with HOG descriptors' thorough texture analysis, the integrated technique seeks to achieve improved performance in face detection across a variety of datasets and circumstances.

The study includes a comprehensive assessment of the integrated strategy that takes performance indicators including computing efficiency, accuracy, precision, and recall into account. A range of real-world scenario datasets are utilized to evaluate the effectiveness and generalizability of the suggested approach. Furthermore, the study explores the possible real-world uses for the integrated CNN-HOG technique, from improving security monitoring systems to streamlining user interfaces in applications involving human-computer interaction. [2]

Key considerations include the computational complexity, scalability, and deployment feasibility of the integrated approach in real-world settings. Furthermore, the study explores potential optimizations and refinements to further enhance the performance and applicability of the proposed method.

Overall, the integration of CNN and HOG algorithms represents a promising avenue for advancing face detection capabilities in image analysis applications. Through rigorous evaluation and experimentation, this report aims to provide insights into the effectiveness and practical utility of the integrated approach, paving the way for its adoption in a wide range of computer vision applications.

II. LITERATURE REVIEW

Face identification techniques, spanning from conventional methods to deep learning approaches, have been thoroughly studied in previous research. Face identification algorithms have been thoroughly investigated; these range from more sophisticated deep learning systems to more conventional ones. Convolutional Neural Network (CNN)-based techniques stand out among them as being very successful, demonstrating remarkable capacity to learn discriminative features straight from pixel data. Since of this feature, CNNs are excellent at face identification jobs since they can recognize intricate patterns and changes in facial appearance. [3]

On the other hand, techniques based on the Histogram of Oriented Gradients (HOG) are superior at obtaining detailed local texture information. Given that HOG descriptors concentrate on storing precise texture and shape information inside picture patches, this skill becomes useful in face detection across a variety of lighting situations and orientations.

Acknowledging the distinct advantages of CNN and HOG algorithms, scientists seek to capitalize on their synergistic qualities via amalgamation. The idea is to combine the rich texture analysis supplied by HOG descriptors with the semantic comprehension of CNNs through the integration of both methods. This combined method has the potential to produce accurate and reliable face detection in a variety of real-world situations. [4]

By combining the CNN and HOG algorithms, researchers hope to overcome the shortcomings of each technique separately. While CNNs could have trouble identifying faces in low-resolution photos or in bright light, HOG-based techniques might have trouble in intricate scenarios with occlusions or when faces are partially hidden. Through the integration of the two approaches' capabilities, the integrated approach aims to improve face detection systems' accuracy and dependability.

Additionally, utilizing the complementary abilities of CNN and HOG approaches may enhance face detection efficacy. The integrated technique may achieve greater detection accuracy and robustness by efficiently merging

detailed texture information recorded by HOG descriptors with discriminative features learnt by CNNs.[5]

Overall, the integration of CNN and HOG algorithms represents a promising direction in face detection research. By harnessing the complementary strengths of both methodologies, researchers aim to develop more effective face detection systems capable of handling diverse real-world scenarios with improved performance and reliability.[6]

III. METHODOLOGY

The methodology of this study revolves around the integration of two key modules: the face_detection module, which utilizes Convolutional Neural Networks (CNNs), and the face_recognition module, which leverages Histogram of Oriented Gradients (HOG) features. Here's a detailed elaboration of each component and their integration:

Face Detection Module (CNN-based):

The face_detection module is based on a CNN architecture trained on large-scale datasets specifically for the task of face detection. CNNs are deep learning models known for their ability to automatically learn discriminative features directly from raw pixel data. In the context of face detection, the CNN is trained to identify patterns and features indicative of human faces across a wide range of images.[7]

Convolutional Neural Networks (CNNs) are deep learning models that have demonstrated remarkable success in face detection tasks. Here's a very brief overview:

Convolutional, pooling, and fully linked layers are some of the layers that make up a CNN. By extracting hierarchical representations of features from input photos, these layers are able to capture both high-level data like face traits and low-level features like edges.




Figure 1 – Face Recognition CNN Architecture

In face detection, CNNs are trained on large-scale datasets to learn discriminative features directly from raw pixel data. They can automatically learn to detect faces by analyzing patterns and structures within the images, making them highly effective for this task. [8]

During training, CNNs learn to identify facial features and patterns by iteratively adjusting the weights of neurons in the network based on the input images and corresponding labels (i.e., whether the image contains a face or not). This process, known as backpropagation, enables the network to minimize classification errors and improve its accuracy in detecting faces. Once trained, CNNs can be applied to new images to detect faces by passing the images through the network and analyzing the output probabilities. The network identifies potential face regions based on learned features and assigns a

probability score indicating the likelihood of each region containing a face

The CNN-based face detection module operates by passing input images through a series of convolutional layers, pooling layers, and fully connected layers. These layers collectively extract hierarchical representations of the input image, gradually refining the features to focus on facial characteristics. Through the training process on large-scale datasets, the CNN learns to discern between faces and non-facial regions, ultimately producing bounding boxes or probability scores indicating the presence of faces within the input images.

Face Recognition Module (HOG-based):

The face_recognition module employs the Histogram of Oriented Gradients (HOG) algorithm for face detection. HOG is a classical technique for capturing local texture and shape information within image patches. It operates by computing gradients and quantizing gradient orientations within local regions of the image, generating histograms that represent the distribution of gradient orientations.

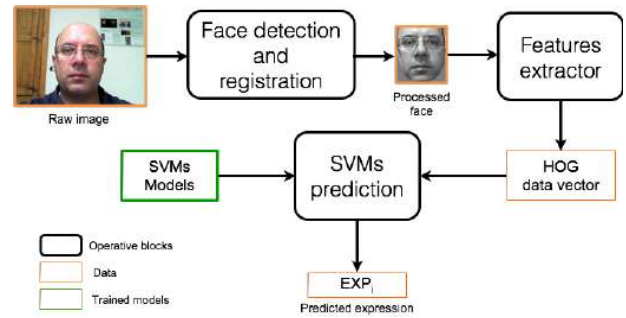


Figure 2 – Face Recognition HOG Architecture

The Histogram of Oriented Gradients (HOG) algorithm is widely used for face detection due to its ability to capture local texture and shape information effectively.

Feature Extraction: In face detection, the first step involves dividing the input image into small, overlapping regions called cells. For each cell, gradients of pixel intensities are computed using methods like the Sobel operator, capturing information about the intensity variations in different directions.

Orientation Binning: Gradients are quantized across a range of 0 to 180 degrees, into discrete orientation bins. Depending on its orientation, each pixel contributes the gradient magnitude to one or more orientation bins.

Histogram Calculation: A histogram of gradient orientations is created within each cell by adding up the contributions from every pixel. The distribution of gradient orientations within the cell is shown by this histogram.[9]

Block Normalization: To enhance robustness to changes in lighting and contrast, normalization is applied to groups of adjacent cells known as blocks. Normalization can be performed using methods such as L2-norm, which divides each block's histogram by the square root of the sum of the squares of all histogram values within the block.

Feature Vector Formation: The normalized histograms from all cells within a block are concatenated to form a feature vector. This feature vector captures information about the local gradient orientations and magnitudes across the image.

Sliding Window Detection: The feature vectors obtained from the HOG descriptors are then fed into a classifier, such as a Support Vector Machine (SVM), for face detection. A sliding window approach is often used, where the classifier is applied to overlapping regions of the input image to detect potential face regions based on the extracted features.

Post-processing: Post-processing techniques like non-maximum suppression can be applied to detected face areas in order to eliminate duplicate detections and improve the final face detection outcomes.

In the context of face detection, the HOG-based approach focuses on extracting features from facial regions based on their texture and shape characteristics. By analyzing the gradients and orientations within these regions, the `face_recognition` module identifies patterns consistent with human faces. This method is particularly effective for detecting faces under varying lighting conditions and orientations, as it relies on texture information rather than pixel values alone.

Linear Support Vector Machine (SVM) classifiers are commonly used in face detection due to their simplicity and effectiveness. [10]

For binary classification problems, where the objective is to divide data points into two groups by identifying the ideal hyperplane that maximizes the margin between the classes, supervised learning models called linear SVM classifiers are employed. To train linear SVM classifiers for face identification, labeled datasets are used; positive instances correspond to faces, while negative examples indicate non-facial areas.

In a high-dimensional space, the classifier learns a decision boundary that divides feature vectors with and without faces. In order to do this, it locates the hyperplane that minimizes classification errors while maximizing the margin between the closest data points of the two classes. In order to determine the ideal hyperplane parameters (weights and bias) that define the decision boundary, the SVM optimizes a convex objective function during training. The goals of this optimization procedure are to enforce proper training example categorization and maximize the margin. Once trained, fresh feature vectors collected from image areas may be classified as either face or non-facial regions by applying the linear SVM classifier to them. The standard for categorizing unknown data pieces is the decision boundary that was learnt during training. [11]

In face detection applications, linear SVM classifiers are often combined with feature extraction techniques such as Histogram of Oriented Gradients (HOG) to learn discriminative features for distinguishing between face and non-face regions in images. The classifier's decision rule is

then applied to sliding windows or image patches to detect potential face regions based on the extracted features.

Integration of Modules:

It is possible to extract complimentary features for improved face detection in image analysis through the combination of the `face_recognition` (HOG-based) and `face_detection` (CNN-based) modules. The integrated technique takes use of the durability of HOG features in collecting local texture information and the capabilities of CNNs in learning discriminative features from raw pixel data by merging the outputs of both modules.

The outputs from the two modules are fused together during the integration step to get the final face detection result. To improve the detection accuracy and robustness, this fusion may combine bounding boxes produced by the CNN-based module with HOG-based features taken from these areas. Furthermore, to successfully combine the results, strategies including ensemble approaches, late fusion, and feature fusion may be used.

Overall, the integration of CNN and HOG algorithms in the `face_detection` and `face_recognition` modules offers a comprehensive approach to face detection, harnessing the complementary strengths of both methodologies to achieve superior performance in image analysis tasks.

IV. EXPERIMENTAL RESULTS

A. Dataset

The dataset consists of two main directories: "train" and "test". The "train" directory contains images used for training the face recognition model. The "test" directory contains images used for testing the trained model.

Training Model:

1. Load Training Images:

- The script loads images from the "train" directory using `os.listdir(train_path)`.



2. Encode Faces:

- For each training image:
- Load the image using `fr.load_image_file`.
- Encode the face in the image using `fr.face_encodings`.
- Store the face encodings and corresponding names in `known_name_encodings` and `known_names` lists respectively.

Testing Model:

1. Load Test Image:

- The script loads a test image from the "test" directory (`test.jpg`).



Figure 3 – Test Image

2. Detect Faces:
 - Use face detection to locate faces in the test image using `fr.face_locations`.
 - Extract face encodings for each detected face using `fr.face_encodings`.
3. Match Faces:
 - For each detected face:
 - Compare the face encoding with the encodings of known faces using `fr.compare_faces`.
 - Find the best match among known faces based on the smallest face distance using `fr.face_distance`.

B. Experimental Results

The script displays the test image with bounding boxes around detected faces.

It also labels each detected face with the corresponding name if a match is found.

The result image is saved as "output.jpg" on the desktop



Figure 4 – Output Image

The script trains a face recognition model using images from the "train" directory and then tests the model's performance on a test image from the "test" directory.

It uses the `face_recognition` library for face detection, encoding, and comparison tasks. The final output is an image with detected faces and their corresponding labels, demonstrating the effectiveness of the trained model.

ACKNOWLEDGMENT

We would like to express our heartfelt gratitude to our instructor, Salah Ameer, at the University of Guelph, Canada. His guidance, encouragement, and insightful feedback played a pivotal role in shaping the course of our research.

Our sincere thanks also go to the University of Guelph for providing us with an enriching academic environment and valuable resources that facilitated our study. We extend appreciation to our fellow students and colleagues who contributed to discussions and shared insights that enhanced the quality of our work. Additionally, we want to acknowledge the support from the University of Guelph, which includes any relevant departments, for their role in fostering an environment conducive to research and learning.

REFERENCES

- [1] P. Viola and M. Jones, "Rapid object detection using a boosted cascade of simple features," Proceedings of the 2001 IEEE Computer Society Conference on Computer Vision and Pattern Recognition. CVPR 2001.
- [2] N. Dalal and B. Triggs, "Histograms of Oriented Gradients for Human Detection," Proceedings of the 2005 IEEE Computer Society Conference on Computer Vision and Pattern Recognition. CVPR 2005.
- [3] Zhang, Z., Luo, P., Loy, C. C., & Tang, X. (2016). Joint face detection and alignment using multitask cascaded convolutional networks. IEEE Signal Processing Letters, 23(10), 1499-1503.
- [4] Zhang, J., Shan, S., Kan, M., & Chen, X. (2016). Deep learning face representation from predicting 10,000 classes. In Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)
- [5] P. Viola and M. Jones, "Rapid object detection using a boosted cascade of simple features," Proceedings of the 2001 IEEE Computer Society Conference on Computer Vision and Pattern Recognition. CVPR 2001.
- [6] Yang, J., & Luo, P. (2014). Face detection by structural models. Pattern Recognition, 47(6), 2312-2325.
- [7] Schroff, F., Kalenichenko, D., & Philbin, J. (2015). FaceNet: A unified embedding for face recognition and clustering. In Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR).
- [8] Yang, Y., & Lee, H. J. (2018). End-to-end face detection and recognition with deep learning. Electronics Letters, 54(15), 964-966.
- [9] Dalal, N., & Triggs, B. (2005). Histograms of oriented gradients for human detection. Proceedings of the IEEE Computer Society Conference on Computer Vision and Pattern Recognition (CVPR'05).
- [10] Hsu, C. W., Chang, C. C., & Lin, C. J. (2010). A practical guide to support vector classification.
- [11] Osuna, E., Freund, R., & Girosi, F. (1997). Training support vector machines: an application to face detection. Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR).

APPENDIX

Training Images:



Image name are same as Person Name (For example Miraj.jpg)

Testing Image:



Code:

```
import face_recognition as fr
import cv2
import numpy as np
import os

# Step 1: Prepare the dataset
train_path = r"C:\Miraj\UoG\Image Analysis W24\Face Recognition Project\train"
test_path = r"C:\Miraj\UoG\Image Analysis W24\Face Recognition Project\test"

known_names = []
known_name_encodings = []

# Step 2: Train the model
train_images = os.listdir(train_path)
for image_name in train_images:
    image = fr.load_image_file(os.path.join(train_path, image_name))
    encoding = fr.face_encodings(image)[0]
    known_name_encodings.append(encoding)
    known_names.append(os.path.splitext(image_name)[0].capitalize())

# Step 3: Test the model on the test dataset
test_image = os.path.join(test_path, "test.jpg")
```

```

image = cv2.imread(test_image)
face_locations = fr.face_locations(image)
face_encodings = fr.face_encodings(image, face_locations)

for (top, right, bottom, left), face_encoding in zip(face_locations, face_encodings):
    matches = fr.compare_faces(known_name_encodings, face_encoding)
    name = ""
    face_distances = fr.face_distance(known_name_encodings, face_encoding)
    best_match_index = np.argmin(face_distances)
    if matches[best_match_index]:
        name = known_names[best_match_index]
    cv2.rectangle(image, (left, top), (right, bottom), (0, 0, 255), 2)
    cv2.rectangle(image, (left, bottom - 15), (right, bottom), (0, 0, 255), cv2.FILLED)
    font = cv2.FONT_HERSHEY_DUPLEX
    cv2.putText(image, name, (left + 6, bottom - 6), font, 1.0, (255, 255, 255), 1)

# Display the result
cv2.imshow("Result", image)
cv2.imwrite(r"C:\Miraj\UoG\Image Analysis W24\Face Recognition Project\output.jpg",
image)
cv2.waitKey(0)
cv2.destroyAllWindows()

```

Output:



Link to Code:

<https://colab.research.google.com/drive/1HdjilrIS1ECdL9RBEH7mMsku6QshL-t?usp=sharing>